Dor Alagem

M.Sc. Computer Science Student · Software Engineer

dor3382@gmail.com 053-3393036 github.com/DoriAlagem
linkedin.com/in/dor-alagem dor-alagem.com

PROFILE

M.Sc. Computer Science student (Machine Learning & Big Data) with a 96 B.Sc. GPA. Builds AI agents and automation pipelines in Python, keeping the LLM's reasoning grounded in deterministic, testable code and backing it with evals, Pytest suites, and CI/CD. Experienced in debugging, REST API validation, and shipping working projects end to end, from data pipeline to live deployment. Strong foundation in operating systems, computer networks, and algorithmic problem solving.

EDUCATION

M.Sc. in Computer Science – Machine Learning & Big Data 10/2026 – Present
The Open University of Israel

B.Sc. in Computer Science 04/2024 – 09/2026
HIT – Holon Institute of Technology

- **GPA: 96**
- **Dean's List** (2024–2025) — Recognized for academic excellence and consistent high performance.

PROJECTS

AI.I – Airport Investment Intelligence Agent
Technologies: Python, FastAPI, Gemini (function calling), SQLite, DuckDB, Pytest, GitHub Actions, Render

- Built a tool-calling AI agent that ranks 396 US airports by suppressed demand for investors: a ReAct-style loop over six tools, grounded in a deterministic scoring engine so the LLM reasons but never produces a figure.
- Engineered an ETL pipeline that condenses ~390 MB of federal aviation data (BTS, FAA; 36 months) into a 7.9 MB snapshot; with no tuning, it ranks Newark and LaGuardia, notoriously congested, first and second among large hubs.
- Developed an agent evaluation harness over a golden dataset that checks every figure in an answer against tool outputs, catching hallucinations that read as fluent prose; each answer ships with its tool-call trace for observability.
- Shipped a live demo with a hands-free English/Hebrew voice mode, delivered through CI/CD: 102 tests on every change and automatic deploys to Render.

AI-Powered Job Scanner
Technologies: Python, Groq LLM API (GPT-OSS 20B), BeautifulSoup, Pytest, GitHub Actions

- Built an autonomous LLM pipeline that scrapes six job sources twice daily, scores each listing against a parsed CV with structured JSON output, and emails a ranked digest; it has screened 11,000+ listings over six months unattended.
- Cut LLM calls 5x by scoring listings in batches, behind a rule-based pre-filter (region, seniority, experience), keeping every run within the free tier; retries and quota handling let a run degrade instead of failing.
- Added persistent deduplication so no listing is scored or sent twice; a new source plugs in with a one-line change.
- Wrote 62 Pytest tests, including an end-to-end run with mocked LLM and SMTP; runs serverless on GitHub Actions. Developed through agentic engineering with Claude Code: context files (CLAUDE.md) and agent-ready issue triage.

SKILLS

Programming Languages: Python, JavaScript, C, SQL

Agentic AI: LLM agents, tool calling, RAG, context engineering, evals, guardrails, prompt engineering

LLMs & AI Tools: Google ADK, Groq, OpenAI APIs, Whisper (speech-to-text), Claude Code

Testing & Automation: Pytest, Playwright, Postman, GitHub Actions (CI)

Backend, Data & Cloud: FastAPI, Node.js, REST APIs, Microservices, SQLite, DuckDB, MongoDB, MySQL, AWS, Render, Git